

Introduction to probabilistic ML

Exercise 1: Multivariate Gaussian

1. Estimate μ_x and Σ_x using MLE.

Proof. We are looking for $\mu_x^{ML}, \Sigma_x^{ML} = \operatorname{argmax}_{\mu_x, \Sigma_x} p(\mathbf{x}|\mu_x, \Sigma_x)$ which is equivalent to $\mu_x^{ML}, \Sigma_x^{ML} = \operatorname{argmax}_{\mu_x, \Sigma_x} \log p(\mathbf{x}|\mu_x, \Sigma_x)$ since the logarithm is monotonic.

Let's calculate the form of $\log p(\mathbf{x}|\mu_x, \Sigma_x)$:

$$\begin{aligned}
 \log p(\mathbf{x}|\mu_x, \Sigma_x) &= \log \prod_{n=1}^N p(x_n|\mu_x, \Sigma_x) = \sum_{n=1}^N \log p(x_n|\mu_x, \Sigma_x) = \\
 &= \sum_{n=1}^N \left[\log \left(\frac{1}{\sqrt{(2\pi)^K |\Sigma_x|}} \right) - \frac{1}{2} (x_n - \mu_x)^\top \Sigma_x^{-1} (x_n - \mu_x) \right] = \\
 &= -\frac{N}{2} \log((2\pi)^K |\Sigma_x|) - \frac{1}{2} \sum_{n=1}^N [(x_n - \mu_x)^\top \Sigma_x^{-1} (x_n - \mu_x)] = \\
 &= C + \frac{N}{2} \log |\Sigma_x^{-1}| - \frac{1}{2} \sum_{n=1}^N [(x_n - \mu_x)^\top \Sigma_x^{-1} (x_n - \mu_x)] = \quad (1) \\
 &= C + \frac{N}{2} \log |\Sigma_x^{-1}| - \frac{1}{2} \sum_{n=1}^N \operatorname{Tr} [(x_n - \mu_x)^\top \Sigma_x^{-1} (x_n - \mu_x)] = \\
 &= C + \frac{N}{2} \log |\Sigma_x^{-1}| - \frac{1}{2} \sum_{n=1}^N \operatorname{Tr} [\Sigma_x^{-1} (x_n - \mu_x) (x_n - \mu_x)^\top] = \\
 &= C + \frac{N}{2} \log |\Sigma_x^{-1}| - \frac{1}{2} \operatorname{Tr} \left[\Sigma_x^{-1} \sum_{n=1}^N (x_n - \mu_x) (x_n - \mu_x)^\top \right] \quad (2)
 \end{aligned}$$

where C is a constant and from equation 1 to 2 we have used three facts: (1) a real number (1×1 matrix) is equal to its trace, (2) $\operatorname{Tr}[AB] = \operatorname{Tr}[BA]$, and (3) the trace is a linear function.

Now let's write the derivative of $\log p(\mathbf{x}|\mu_x, \Sigma_x)$ w.r.t. μ_x using the expression 1.

$$\begin{aligned} \frac{\partial}{\partial \mu_x} \log p(\mathbf{x}|\mu_x, \Sigma_x) &= -\frac{1}{2} \sum_{n=1}^N \frac{\partial}{\partial \mu_x} [(x_n - \mu_x)^\top \Sigma_x^{-1} (x_n - \mu_x)] = \\ &= -\frac{1}{2} \sum_{n=1}^N [-\Sigma_x^{-1} (x_n - \mu_x) - (x_n - \mu_x)^\top \Sigma_x^{-1}] = \\ &= -\frac{1}{2} \sum_{n=1}^N [-2\Sigma_x^{-1} (x_n - \mu_x)] = \Sigma_x^{-1} \sum_{n=1}^N (x_n - \mu_x) \end{aligned}$$

where we have used the fact that Σ_x^{-1} is a full-rank symmetric positive-definite matrix.

And we do the same w.r.t. Σ_x^{-1} using the expression 2.

$$\frac{\partial}{\partial \Sigma_x^{-1}} \log p(\mathbf{x}|\mu_x, \Sigma_x) = \underbrace{\frac{N}{2} \frac{\partial \log |\Sigma_x^{-1}|}{\partial \Sigma_x^{-1}}}_{(3)} - \underbrace{\frac{1}{2} \frac{\partial \text{Tr}}{\partial \Sigma_x^{-1}} \left[\Sigma_x^{-1} \sum_{n=1}^N (x_n - \mu_x)(x_n - \mu_x)^\top \right]}_{(4)} = (*)$$

$$(3) = \frac{N}{2} (\Sigma_x^{-1})^{-\top} = \frac{N}{2} \Sigma_x^\top \quad \text{since } \frac{\partial \log |A|}{\partial A} = A^{-\top}$$

$$\begin{aligned} (4) &= \frac{1}{2} \left[\sum_{n=1}^N (x_n - \mu_x)(x_n - \mu_x)^\top \right]^\top \quad \text{since } \frac{\partial \text{Tr}(AB)}{\partial A} = B^\top \\ &= \frac{1}{2} \sum_{n=1}^N (x_n - \mu_x)(x_n - \mu_x)^\top \quad \text{since } (A+B)^\top = A^\top + B^\top \text{ and } (AB)^\top = B^\top A^\top \end{aligned}$$

Thus

$$(*) = (3) - (4) = \frac{N}{2} \Sigma_x^\top - \frac{1}{2} \sum_{n=1}^N (x_n - \mu_x)(x_n - \mu_x)^\top$$

Therefore we have the equation system:

$$\begin{cases} \partial_{\mu_x} \log p(\mathbf{x}) = 0 & \iff \Sigma_x^{-1} \sum_{n=1}^N (x_n - \mu_x) = 0 \\ \partial_{\Sigma_x^{-1}} \log p(\mathbf{x}) = 0 & \iff \frac{N}{2} \Sigma_x^\top - \frac{1}{2} \sum_{n=1}^N (x_n - \mu_x)(x_n - \mu_x)^\top = 0 \end{cases}$$

The first equation can be readily solved since

$$\Sigma_x^{-1} \sum_{n=1}^N (x_n - \mu_x) = 0 \iff \sum_{n=1}^N x_n - N\mu_x = 0 \iff \mu_x = \frac{1}{N} \sum_{n=1}^N x_n \quad (3)$$

and we can check that it is in fact a maximum

$$\frac{\partial^2}{\partial \mu_x} \log p(\mathbf{x}|\mu_x, \Sigma_x) = -N\Sigma_x^{-1} \prec 0 \quad \text{since } \Sigma_x \succ 0 \text{ (p.s.d.)} \quad (4)$$

so we can call $\mu_x^{ML} := \frac{1}{N} \sum_{n=1}^N x_n$ and substituting μ_x^{ML} in the second equation we have that

$$\begin{aligned} \frac{N}{2} \Sigma_x^\top - \frac{1}{2} \sum_{n=1}^N (x_n - \mu_x^{ML})(x_n - \mu_x^{ML})^\top &= 0 \iff \\ \frac{N}{2} \Sigma_x^\top &= \frac{1}{2} \sum_{n=1}^N (x_n - \mu_x^{ML})(x_n - \mu_x^{ML})^\top \iff \\ \Sigma_x^\top &= \Sigma_x = \frac{1}{N} \sum_{n=1}^N (x_n - \mu_x^{ML})(x_n - \mu_x^{ML})^\top \end{aligned}$$

and again we can check that this is a maximum:

$$\frac{\partial^2}{\partial \Sigma_x^{-1}} \log(\mathbf{x}|\mu_x, \Sigma_x) = \frac{N}{2} \frac{\partial}{\partial \Sigma_x^{-1}} (\Sigma_x^{-1})^{-1} = -\frac{N}{2} \Sigma_x^2 \prec 0 \quad (5)$$

Finally, $\mu_x^{ML} := \frac{1}{N} \sum_{n=1}^N x_n$ and $\Sigma_x^{ML} := \frac{1}{N} \sum_{n=1}^N (x_n - \mu_x^{ML})(x_n - \mu_x^{ML})^\top$. \square

2. Estimate μ_x using MAP assuming known μ_x and prior $\mu_x \sim \mathcal{N}(\mu_0, \Sigma_0)$.

Proof. Using Bayes' theorem:

$$p(\mu_x|\mathbf{x}, \mu_0, \Sigma_0, \Sigma_x) = \frac{p(\mathbf{x}|\mu_x, \Sigma_x)p(\mu_x|\mu_0, \Sigma_0)}{p(\mathbf{x}|\mu_0, \Sigma_0, \Sigma_x)} \propto p(\mathbf{x}|\mu_x, \Sigma_x)p(\mu_x|\mu_0, \Sigma_0) \quad (1)$$

We are going to discover the form of the posterior distribution by trying to obtain a formula that we can recognize, if we achieve this calculating the normalizing constant is straight-forward. In particular, we are going to compute $\log p(\mu_x|\mathbf{x}, \mu_0, \Sigma_0, \Sigma_x)$ and try to obtain a quadratic form of μ_x which is the form that gaussian distributions have.

$$\begin{aligned} \log p(\mu_x|\mathbf{x}, \mu_0, \Sigma_0, \Sigma_x) &= \log \mathcal{N}(\mathbf{x}|\mu_x, \Sigma_x) + \log \mathcal{N}(\mu_x|\mu_0, \Sigma_0) + C = \\ &= -\frac{1}{2} \sum_{n=1}^N (x_n - \mu_x)^\top \Sigma_x^{-1} (x_n - \mu_x) - \frac{1}{2} (\mu_x - \mu_0)^\top \Sigma_0^{-1} (\mu_x - \mu_0) + C = \\ &= -\frac{1}{2} \left[\sum_{n=1}^N (\mu_x^\top \Sigma_x^{-1} \mu_x - 2\mu_x^\top \Sigma_x^{-1} x_n) + \mu_x^\top \Sigma_0^{-1} \mu_x - 2\mu_x^\top \Sigma_0^{-1} \mu_0 \right] + C = \\ &= -\frac{1}{2} \left[\mu_x^\top (N\Sigma_x^{-1} + \Sigma_0^{-1}) \mu_x - 2\mu_x^\top \left(\Sigma_x^{-1} \sum_{n=1}^N x_n + \Sigma_0^{-1} \mu_0 \right) \right] + C \end{aligned} \quad (2)$$

Now, we have to complete squares in equation 2. To do that we know that, if A is symmetric, $(x - y)^\top A(x - y) = x^\top Ax + y^\top Ay - 2x^\top Ay$. Comparing equation 2 with the previous formula we can call $x = \mu_x$ and $A = (N\Sigma_x^{-1} + \Sigma_0^{-1})$.

In order to find out who is y we have to make A appear in the expression $-2x^\top Ay$ of equation 2. We can easily achieve this by multiplying by AA^{-1} , making equation 2 to look like

$$(2) = -\frac{1}{2} \left[\mu_x^\top A \mu_x - 2\mu_x^\top A \left[A^{-1} \left(\Sigma_x^{-1} \sum_{n=1}^N x_n + \Sigma_0^{-1} \mu_0 \right) \right] \right] + C$$

and by calling $y = A^{-1} \left(\Sigma_x^{-1} \sum_{n=1}^N x_n + \Sigma_0^{-1} \mu_0 \right)$ we have that

$$(2) = -\frac{1}{2} (\mu_x - y)^\top A (\mu_x - y) + C$$

Now, if μ_x had a normal posterior distribution, i.e., $\mu_x | \mathbf{x} \sim \mathcal{N}(\mu_1, \Sigma_1)$, then $\log p(\mu_x | x)$ would be of the form

$$\log p(\mu_x | \mathbf{x}) = -\frac{1}{2} (\mu_x - \mu_1)^\top \Sigma_1^{-1} (\mu_x - \mu_1) + C$$

which implies, by comparing the two expressions, that the posterior distribution of μ_x is a Gaussian distribution with mean $\mu_1 = y$ and covariance $\Sigma_1 = A^{-1}$.

Finally, we need to compute the MAP estimate of μ_x given \mathbf{x} . This estimator is defined as $\mu_x^{MAP} := \operatorname{argmax}_{\mu_x} p(\mu_x | \mathbf{x})$ which, making similar calculations as the ones done in the previous section, can be proved to be the mean of the normal distribution, that is, $\mu_x^{MAP} = \mu_1 = y$. \square

Exercise 2: Categorical distribution

1. Estimate π using ML.

Proof. We have to solve the problem (note that we use the shorthand $\pi = \{\pi_k\}_{k=1}^K$)

$$\pi^{ML} := \operatorname{argmax}_{\pi} p(\mathbf{x} | \pi) \quad \text{subject to} \quad \sum_{k=1}^K \pi_k = 1 \quad (1)$$

which is equivalent to solving

$$\pi^{ML} := \operatorname{argmax}_{\pi} \log p(\mathbf{x} | \pi) \quad \text{subject to} \quad \sum_{k=1}^K \pi_k = 1 \quad (2)$$

and using Lagrange multipliers this is equivalent to solving

$$\pi^{ML} := \operatorname{argmax}_{\pi} \left[\log p(\mathbf{x}|\pi) - \lambda \left(\sum_{k=1}^K \pi_k - 1 \right) \right] \quad (3)$$

where λ is a sufficiently large real positive number.

Let's write down the form of the log-likelihood:

$$\begin{aligned} p(\mathbf{x}|\pi) &= \prod_{n=1}^N p(x_n|\pi) = \prod_{n=1}^N \prod_{k=1}^K \pi_k^{[x_n=k]} \quad \text{where } [x=k] = \begin{cases} 1 & \text{if } x=k \\ 0 & \text{otherwise} \end{cases} \\ \log p(\mathbf{x}|\pi) &= \sum_{n=1}^N \sum_{k=1}^K \log \left(\pi_k^{[x_n=k]} \right) = \sum_{n=1}^N \sum_{k=1}^K [x_n=k] \log \pi_k \end{aligned} \quad (4)$$

Now we have to solve the system

$$\begin{cases} \frac{\partial}{\partial \pi_1} \left[\log p(\mathbf{x}|\pi) - \lambda \left(\sum_{k=1}^K \pi_k - 1 \right) \right] = 0 \\ \frac{\partial}{\partial \pi_2} \left[\log p(\mathbf{x}|\pi) - \lambda \left(\sum_{k=1}^K \pi_k - 1 \right) \right] = 0 \\ \dots \\ \frac{\partial}{\partial \pi_K} \left[\log p(\mathbf{x}|\pi) - \lambda \left(\sum_{k=1}^K \pi_k - 1 \right) \right] = 0 \end{cases} \quad (5)$$

Therefore, let us solve this equation for every $l \in \{1, 2, \dots, K\}$

$$\begin{aligned} \frac{\partial \log p(\mathbf{x}|\pi)}{\partial \pi_l} &= \sum_{n=1}^N \sum_{k=1}^K \frac{\partial ([x_n=k] \log \pi_k)}{\partial \pi_l} - \lambda \frac{\partial \left(\sum_{k=1}^K \pi_k - 1 \right)}{\partial \pi_l} = \\ &= \sum_{n=1}^N \frac{[x_n=l]}{\pi_l} - \lambda = 0 \iff \pi_l = \frac{1}{\lambda} \sum_{n=1}^N [x_n=l] = \frac{1}{\lambda} n_l \end{aligned}$$

where n_l represents how many x_n in \mathbf{x} have the value l . Note that this is indeed a maximum since

$$\frac{\partial^2}{\partial \pi_l^2} \log p(\mathbf{x}|\pi) = -\frac{n_l}{\pi_l^2} < 0$$

assuming that every class has a non-zero probability of happening (that is, it has been observed at least once).

We have a set of solutions $\pi_k^{ML}(\lambda) = n_k/\lambda$, one per each value of λ . In order to solve the problem we solve λ substituting $\pi^{ML}(\lambda)$ on the restriction over π :

$$\sum_{k=1}^K \pi_k^{ML}(\lambda) = \frac{1}{\lambda} \sum_{k=1}^K n_k = 1 \iff \lambda = \sum_{k=1}^K n_k = N \quad (6)$$

Therefore, the maximum likelihood estimator of π_k is

$$\pi_k^{ML} = \frac{1}{N} \sum_{n=1}^N [x_n = k] = \frac{n_k}{N} \quad (7)$$

□

2. Calculate the posterior $p(\pi|\mathbf{x})$ assuming a prior $\pi \sim \text{Dirichlet}(\alpha)$.

Proof. We assume a prior

$$p(\pi|\alpha) = \text{Dirichlet}(\alpha) = \frac{1}{B(\alpha)} \prod_{k=1}^K \pi_k^{\alpha_k - 1} \quad (1)$$

Using Bayes' theorem we have that

$$\begin{aligned} p(\pi|\mathbf{x}, \alpha) &\propto p(\mathbf{x}|\pi)p(\pi|\alpha) \propto \prod_{n=1}^N \prod_{k=1}^K \pi_k^{[x_n=k]} \prod_{k=1}^K \pi_k^{\alpha_k - 1} = \\ &= \prod_{k=1}^K \pi_k^{\sum_{n=1}^N [x_n=k] + \alpha_k - 1} = \prod_{k=1}^K \pi_k^{n_k + \alpha_k - 1} \end{aligned}$$

And, since it has the same form as a Dirichet distribution up to the normalization constant, we know that $\pi|\mathbf{x} \sim \text{Dirichlet}(n_1 + \alpha_1, n_2 + \alpha_2, \dots, n_K + \alpha_K)$. □

Exercise 3: Graphical models

1. Write the graphical model corresponding to the generative model

$$p(\{x_n, z_n\}_{n=1}^N, \{\pi_k, \mu_k\}_{k=1}^K) = \prod_{n=1}^N p(x_n | z_n, \{\mu_k\}_{k=1}^K, \sigma_x) p(z_n | \{\pi_k\}_{k=1}^K) p(\{\pi_k\}_{k=1}^K | \alpha)$$




Figure 1: Graphical model of exercise 3.1

2. Write down the generative model of the graphical model.

$$p(\{\omega_n, z_n\}_{n=1}^N, \{\theta_m\}_{m=1}^M, \alpha, \beta) = \prod_{n=1}^N p(\omega_n | z_n, \beta) p(z_n | \{\theta_m\}_{m=1}^M) p(\{\theta_m\}_{m=1}^M | \alpha) p(\beta)$$